

# Caching in ASP.NET Core

## Using LRU Cache (Least Recently Used) and LFU Cache (Least Frequently Used)

## Introduction to Caching

**Caching** is a technique that stores frequently accessed data in a **fast-access storage layer** to reduce latency and improve performance.

## Why Cache?

Without Cache: User Request → Database (Slow) → Response (200-500ms)

# ⌚ LRU Cache (Least Recently Used)

## ⌚ How LRU Cache Works

**Concept:** "If you haven't used it recently, you probably won't need it soon"

Data Structure:

- Hash Map (Dictionary) for  $O(1)$  lookups
- Doubly Linked List for  $O(1)$  reordering

Operations:

1. GET(key): Move item to front (most recent)
2. PUT(key, value):
  - If exists: Update value, move to front
  - If full: Remove tail (least recent), add to front

## ⌚ When to Use LRU:

- ✓ **Browser Cache** - Caching web pages and resources
- ✓ **Operating System Page Cache** - Memory management
- ✓ **Database Query Cache** - Frequently executed queries
- ✓ **Session Management** - User session data
- ✓ **API Response Caching** - REST API responses

# LFU Cache (Least Frequently Used)

## How LFU Cache Works

**Concept:** "If you haven't used it much, you probably don't need it"

Data Structure:

1. Hash Map: key → (value, frequency, node)
2. Frequency Map: frequency → Doubly Linked List
3. Min Frequency tracker

Operations:

1. GET(key): Increase frequency, move to next freq list
2. PUT(key, value):
  - If exists: Update value, increase frequency
  - If full: Remove from min frequency list

## When to Use LFU:

- ✓ **Content Delivery Networks (CDN)** - Popular static content
- ✓ **Search Engine Results** - Frequently searched queries
- ✓ **E-commerce Product Catalog** - Hot products
- ✓ **Advertisement Serving** - Frequently displayed ads
- ✓ **Compiler Optimization** - Frequently used code paths

# LRU vs LFU Comparison

Choose LRU When:	Choose LFU When:
<ul style="list-style-type: none"><li>Recent access matters more than frequency</li></ul>	<ul style="list-style-type: none"><li>Frequency of access matters more than recency</li></ul>
<ul style="list-style-type: none"><li>Sequential access patterns exist</li></ul>	<ul style="list-style-type: none"><li>Stable access patterns with identifiable popular items</li></ul>
<ul style="list-style-type: none"><li>Memory is constrained</li></ul>	<ul style="list-style-type: none"><li>Can tolerate additional memory overhead</li></ul>
<ul style="list-style-type: none"><li>Simple implementation is needed</li></ul>	<ul style="list-style-type: none"><li>Popular content needs to stay cached longer</li></ul>
<ul style="list-style-type: none"><li>Cache pollution from one-time accesses is a concern</li></ul>	<ul style="list-style-type: none"><li>Long-term popularity is important</li></ul>

## ☞ Key Takeaways



### Performance Impact

Caching can reduce response times from 200-500ms to 5-20ms.



### Implementation Complexity

LRU is simpler to implement than LFU, which requires frequency tracking.



### Temporal Locality

LRU exploits temporal locality (recently used items will likely be used again).



### Frequency Patterns

LFU exploits frequency patterns (frequently used items will likely be used again).

## Implementation Summary

### LRU Implementation Tips:

- Use Dictionary + LinkedList in C#
- Implement cache size limit
- Update recency on every access
- Consider thread safety for web applications

### LFU Implementation Tips:

- Use Dictionary + Dictionary of LinkedLists
- Track minimum frequency
- Handle frequency promotion carefully
- Watch for memory overhead

---

Suraj Goud